

Ex3 - Getting and Knowing your Data

This time we are going to pull data directly from the internet. Special thanks to: <https://github.com/justmarkham> for sharing the dataset and materials.

Step 1. Import the necessary libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
```

Step 2. Import the dataset from this [address](https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u.user).

```
In [2]: 'https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u.user'
```

```
Out[2]: 'https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u.user'
```

Step 3. Assign it to a variable called users and use the 'user_id' as index

```
In [3]: users=pd.read_csv('https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u.
```

Step 4. See the first 25 entries

```
In [4]: users.head()
```

```
Out[4]:
```

	age	gender	occupation	zip_code
user_id				
1	24	M	technician	85711
2	53	F	other	94043
3	23	M	writer	32067
4	24	M	technician	43537
5	33	F	other	15213

Step 5. See the last 10 entries

```
In [5]: users.tail(10)
```

```
Out[5]:
```

	age	gender	occupation	zip_code
user_id				
934	61	M	engineer	22902
935	42	M	doctor	66221
936	24	M	other	32789

	age	gender	occupation	zip_code
user_id				
937	48	M	educator	98072
938	38	F	technician	55038
939	26	F	student	33319
940	32	M	administrator	02215
941	20	M	student	97229
942	48	F	librarian	78209
943	22	M	student	77841

Step 6. What is the number of observations in the dataset?

```
In [6]: users.shape[0]
```

```
Out[6]: 943
```

Step 7. What is the number of columns in the dataset?

```
In [7]: users.shape[1]
```

```
Out[7]: 4
```

Step 8. Print the name of all the columns.

```
In [8]: users.columns
```

```
Out[8]: Index(['age', 'gender', 'occupation', 'zip_code'], dtype='object')
```

Step 9. How is the dataset indexed?

```
In [9]: users.index
```

```
Out[9]: Int64Index([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10,
                  ...,
                  934, 935, 936, 937, 938, 939, 940, 941, 942, 943],
                  dtype='int64', name='user_id', length=943)
```

Step 10. What is the data type of each column?

```
In [10]: users.dtypes
```

```
Out[10]: age          int64
gender         object
occupation     object
zip_code       object
dtype: object
```

Step 11. Print only the occupation column

```
In [11]: users.occupation
```

```
Out[11]: user_id
1      technician
2      other
3      writer
4      technician
5      other
...
939     student
940 administrator
941     student
942     librarian
943     student
Name: occupation, Length: 943, dtype: object
```

```
In [12]: users['occupation']

#this is also how we can show the results.
```

```
Out[12]: user_id
1      technician
2      other
3      writer
4      technician
5      other
...
939     student
940 administrator
941     student
942     librarian
943     student
Name: occupation, Length: 943, dtype: object
```

Step 12. How many different occupations are in this dataset?

```
In [13]: users.occupation.nunique()
```

```
Out[13]: 21
```

Step 13. What is the most frequent occupation?

```
In [14]: users.occupation.value_counts().head(1).index[0]
```

```
Out[14]: 'student'
```

Step 14. Summarize the DataFrame.

```
In [15]: users.describe()
```

```
Out[15]:
```

	age
count	943.000000
mean	34.051962
std	12.192740
min	7.000000

	age
25%	25.000000
50%	31.000000
75%	43.000000
max	73.000000

Step 15. Summarize all the columns

In [16]: `users.describe(include='all')`

Out[16]:

	age	gender	occupation	zip_code
count	943.000000	943	943	943
unique	NaN	2	21	795
top	NaN	M	student	55414
freq	NaN	670	196	9
mean	34.051962	NaN	NaN	NaN
std	12.192740	NaN	NaN	NaN
min	7.000000	NaN	NaN	NaN
25%	25.000000	NaN	NaN	NaN
50%	31.000000	NaN	NaN	NaN
75%	43.000000	NaN	NaN	NaN
max	73.000000	NaN	NaN	NaN

Step 16. Summarize only the occupation column

In [18]: `users.occupation.describe()`

Out[18]:

```
count      943
unique      21
top      student
freq      196
Name: occupation, dtype: object
```

16.a Summarize only the age column

In [19]: `users.age.describe()`

Out[19]:

```
count      943.000000
mean       34.051962
std        12.192740
min         7.000000
25%        25.000000
50%        31.000000
75%        43.000000
max        73.000000
Name: age, dtype: float64
```

16.b Summarize only the gender column

```
In [20]: users.gender.describe()
```

```
Out[20]: count      943  
unique        2  
top           M  
freq         670  
Name: gender, dtype: object
```

16.c Summarize only the zip_code column

```
In [21]: users.zip_code.describe()
```

```
Out[21]: count      943  
unique     795  
top      55414  
freq         9  
Name: zip_code, dtype: object
```

Step 17. What is the mean age of users?

```
In [22]: round(users.age.mean())
```

```
Out[22]: 34
```

```
In [28]: round(users.age.std())
```

```
Out[28]: 12
```

```
In [29]: round(users.age.min())
```

```
Out[29]: 7
```

```
In [30]: round(users.age.max())
```

```
Out[30]: 73
```

```
In [40]: round(users.age.median())
```

```
Out[40]: 31
```

```
In [41]: round(users.age.mode())
```

```
Out[41]: 0      30  
dtype: int64
```

```
In [42]: round(users.age.kurtosis())
```

```
Out[42]: 0
```

```
In [46]:
```

```
round(users.age.skew())
```

Out[46]: 1

Step 18. What is the age with least occurrence?

In [55]:

```
users.age.value_counts().tail(10)
```

Out[55]:

14	3
62	2
69	2
64	2
68	2
7	1
66	1
10	1
11	1
73	1

Name: age, dtype: int64